

Linear Classifiers

CS114B Lab 2

Kenneth Lai

January 27, 2023

Vectors and Vector Spaces

- ▶ Wikipedia: A **vector space** is a set of objects called **vectors**, which may be added together and multiplied by numbers called **scalars**

Vectors and Vector Spaces

- ▶ Wikipedia: A **vector space** is a set of objects called **vectors**, which may be added together and multiplied by numbers called **scalars**
- ▶ Lots of examples of vector spaces; we will focus on a particular type of real vector space/coordinate space, where
 - ▶ Vectors are n -tuples of real numbers, for some natural number n
 - ▶ Scalars are real numbers

Vectors and Vector Spaces

- ▶ Wikipedia: A **vector space** is a set of objects called **vectors**, which may be added together and multiplied by numbers called **scalars**
- ▶ Lots of examples of vector spaces; we will focus on a particular type of real vector space/coordinate space, where
 - ▶ Vectors are elements of \mathbb{R}^n
 - ▶ Scalars are elements of \mathbb{R}

Vectors and Vector Spaces

- ▶ Wikipedia: A **vector space** is a set of objects called **vectors**, which may be added together and multiplied by numbers called **scalars**
- ▶ Lots of examples of vector spaces; we will focus on a particular type of real vector space/coordinate space, where
 - ▶ Vectors are elements of \mathbb{R}^n
 - ▶ Scalars are elements of \mathbb{R}
- ▶ The **dimension** of such a vector space (not to be confused with a dimension, i.e., axis, of a Numpy array) is n

Vectors in NLP

- ▶ Idea: objects of interest (e.g., documents, words, etc.) can be represented as vectors (**embeddings**) in some vector space
- ▶ Coordinates of the vector correspond to **features** of the object

Vectors in NLP

- ▶ Idea: objects of interest (e.g., documents, words, etc.) can be represented as vectors (**embeddings**) in some vector space
- ▶ Coordinates of the vector correspond to **features** of the object
 - ▶ Sometimes, these features are human-interpretable
 - ▶ Naïve Bayes features: word counts in a document

Vectors in NLP

- ▶ Idea: objects of interest (e.g., documents, words, etc.) can be represented as vectors (**embeddings**) in some vector space
- ▶ Coordinates of the vector correspond to **features** of the object
 - ▶ Sometimes, these features are human-interpretable
 - ▶ Naïve Bayes features: word counts in a document
 - ▶ Sometimes, they are not
 - ▶ Many word vector “features”

Linear Classifiers

- ▶ Suppose we have a classification problem (e.g., text classification, sequence labeling, etc.)

Linear Classifiers

- ▶ Suppose we have a classification problem (e.g., text classification, sequence labeling, etc.)
- ▶ **Linear classifiers** make their classification decisions based on a **linear combination** of features
 - ▶ Logistic regression
 - ▶ Perceptron
 - ▶ Naïve Bayes (in a way)
 - ▶ ...

Linear Classifiers

- ▶ Let $\mathbf{x} = [x_1 \ \dots \ x_n]$ be a feature vector, $\theta = [\theta_1 \ \dots \ \theta_n]$ be a vector of parameters, g be a classification function, \hat{y} be the classification decision, and \cdot denote the dot product

$$\hat{y} = g\left(\sum_{j=1}^n \theta_j x_j\right) = g(\theta \cdot \mathbf{x})$$

Linear Classifiers

- ▶ Sometimes (especially in logistic regression), within the set of parameters, we distinguish between weights w_j and a bias term b
 - ▶ This is equivalent to having a “dummy feature” with value 1

$$\theta = [w_1 \quad \dots \quad w_n \quad b] = [\mathbf{w} \mid b]$$

$$\mathbf{x}' = [x_1 \quad \dots \quad x_n \quad 1] = [\mathbf{x} \mid 1]$$

$$\hat{y} = g\left(\sum_{j=1}^n w_j x_j + b\right) = g(\mathbf{w} \cdot \mathbf{x} + b) = g(\theta \cdot \mathbf{x}')$$

Linear Classifiers

- ▶ What is g ?
 - ▶ A dot product of two vectors produces a scalar, but in general, we don't just want an arbitrary real number
 - ▶ Sometimes, we want a probability (logistic regression)
 - ▶ Sometimes, we just want the decision itself (perceptron)

Linear Classifiers

- ▶ Let $z = \theta \cdot \mathbf{x}$ (or $\mathbf{w} \cdot \mathbf{x} + b$)
- ▶ Logistic regression: logistic (sigmoid) function
 - ▶ $\sigma(z) = \frac{1}{1 + e^{-z}}$


Linear Classifiers

- ▶ Let $z = \theta \cdot \mathbf{x}$ (or $\mathbf{w} \cdot \mathbf{x} + b$)
- ▶ Logistic regression: logistic (sigmoid) function
 - ▶ $\sigma(z) = \frac{1}{1 + e^{-z}}$
- ▶ Perceptron: (Heaviside) step function
 - ▶ $H(z) = \begin{cases} 1, & \text{if } z > 0 \\ 0, & \text{if } z < 0 \end{cases}$
 - ▶ Sometimes you may see values 1 and -1 , instead of 1 and 0

Linear Classifiers


- ▶ Let $z = \theta \cdot \mathbf{x}$ (or $\mathbf{w} \cdot \mathbf{x} + b$)
- ▶ Logistic regression: logistic (sigmoid) function
 - ▶ $\sigma(z) = \frac{1}{1 + e^{-z}}$
- ▶ Perceptron: (Heaviside) step function
 - ▶ $H(z) = \begin{cases} 1, & \text{if } z > 0 \\ 0, & \text{if } z < 0 \end{cases}$
 - ▶ Sometimes you may see values 1 and -1 , instead of 1 and 0
 - ▶ What if $z = 0$?
 - ▶ Set by convention (1, 0, or $1/2$)

Graphical Representation of a Linear Classifier (1)




Graphical Representation of a Linear Classifier (1)

- Input (including dummy feature 1)




Graphical Representation of a Linear Classifier (1)

- Parameters (weights and bias term)




Graphical Representation of a Linear Classifier (1)

► Sum function Σ




Graphical Representation of a Linear Classifier (1)



► “Score”

- Sometimes (if g is the logistic function) called a **logit**


Graphical Representation of a Linear Classifier (1)




- Classification function g
 - Logistic, step, etc.
 - Later called an activation function

Graphical Representation of a Linear Classifier (1)


► Output



Graphical Representation of a Linear Classifier (1)




Graphical Representation of a Linear Classifier (1)



Graphical Representation of a Linear Classifier (1)

► Output


► Input



Graphical Representation of a Linear Classifier (2)

► Output

► Input




Graphical Representation of a Linear Classifier (2)

► Output

► Input

► $\hat{y} = g(\theta \cdot \mathbf{x})$

► We will assume that the dummy feature 1 is part of \mathbf{x}



Multi-Class Classification

- ▶ Two-class (binary) classification
 - ▶ Compute “score” $z = \theta \cdot \mathbf{x}$ (or $\mathbf{w} \cdot \mathbf{x} + b$)
 - ▶ Compute decision \hat{y} as a function of z
 - ▶ If \hat{y} is interpreted as the probability of (or indicator for) one class, $1 - \hat{y}$ is the probability of (indicator for) the other class

Multi-Class Classification

- ▶ Two-class (binary) classification
 - ▶ Compute “score” $z = \theta \cdot \mathbf{x}$ (or $\mathbf{w} \cdot \mathbf{x} + b$)
 - ▶ Compute decision \hat{y} as a function of z
 - ▶ If \hat{y} is interpreted as the probability of (or indicator for) one class, $1 - \hat{y}$ is the probability of (indicator for) the other class
- ▶ Multi-class (multinomial) classification
 - ▶ Compute a **vector** of scores $\mathbf{z} = \mathbf{\Theta}\mathbf{x}$ (or $\mathbf{W}\mathbf{x} + \mathbf{b}$)
 - ▶ Compute decision \hat{y} as a function of \mathbf{z}

Matrices and Linear Maps

- ▶ A **matrix** is a rectangular array of scalars
- ▶ Two uses of matrices
 1. Matrices represent **linear maps**
 - ▶ Transformations between vector spaces
 - ▶ Given a feature vector \mathbf{x} , we want a score vector \mathbf{z}
 - ▶ $\mathbf{z} = \mathbf{\Theta}\mathbf{x}$ (or $\mathbf{W}\mathbf{x} + \mathbf{b}$)

Matrices and Linear Maps

- ▶ A **matrix** is a rectangular array of scalars
- ▶ Two uses of matrices
 1. Matrices represent **linear maps**
 - ▶ Transformations between vector spaces
 - ▶ Given a feature vector \mathbf{x} , we want a score vector \mathbf{z}
 - ▶ $\mathbf{z} = \mathbf{\Theta}\mathbf{x}$ (or $\mathbf{W}\mathbf{x} + \mathbf{b}$)
 2. Matrices represent **data**
 - ▶ Stacks of feature vectors
 - ▶ Matrix-vector products become matrix-matrix products

Matrices and Linear Maps

- ▶ Warning! A note on notation:
 - ▶ Math convention: A p -by- n matrix defines a linear map from \mathbb{R}^n to \mathbb{R}^p
 - ▶ Matrices have shapes (output dimension, input dimension)
 - ▶ Let $\Theta \in \mathbb{R}^{p \times n}$, $\mathbf{x} \in \mathbb{R}^n$, and $\mathbf{z} \in \mathbb{R}^p$
 - ▶ $\mathbf{z} = \Theta \mathbf{x}$ (or $\mathbf{W}\mathbf{x} + \mathbf{b}$)

Matrices and Linear Maps

- ▶ Warning! A note on notation:
 - ▶ Math convention: A p -by- n matrix defines a linear map from \mathbb{R}^n to \mathbb{R}^p
 - ▶ Matrices have shapes (output dimension, input dimension)
 - ▶ Let $\Theta \in \mathbb{R}^{p \times n}$, $\mathbf{x} \in \mathbb{R}^n$, and $\mathbf{z} \in \mathbb{R}^p$
 - ▶ $\mathbf{z} = \Theta \mathbf{x}$ (or $\mathbf{W} \mathbf{x} + \mathbf{b}$)
 - ▶ Computer science convention (mostly): An n -by- p matrix defines a linear map from \mathbb{R}^n to \mathbb{R}^p (technically $\mathbb{R}^{1 \times n}$ to $\mathbb{R}^{1 \times p}$)
 - ▶ Matrices have shapes (input dimension, output dimension)
 - ▶ Let $\Theta \in \mathbb{R}^{n \times p}$, $\mathbf{x} \in \mathbb{R}^{1 \times n}$, and $\mathbf{z} \in \mathbb{R}^{1 \times p}$
 - ▶ $\mathbf{z} = \mathbf{x} \Theta$ (or $\mathbf{x} \mathbf{W} + \mathbf{b}$)

Matrices and Linear Maps

- ▶ Warning! A note on notation:
 - ▶ Math convention: A p -by- n matrix defines a linear map from \mathbb{R}^n to \mathbb{R}^p
 - ▶ Matrices have shapes (output dimension, input dimension)
 - ▶ Let $\Theta \in \mathbb{R}^{p \times n}$, $\mathbf{x} \in \mathbb{R}^n$, and $\mathbf{z} \in \mathbb{R}^p$
 - ▶ $\mathbf{z} = \Theta \mathbf{x}$ (or $\mathbf{W} \mathbf{x} + \mathbf{b}$)
 - ▶ Computer science convention (mostly): An n -by- p matrix defines a linear map from \mathbb{R}^n to \mathbb{R}^p (technically $\mathbb{R}^{1 \times n}$ to $\mathbb{R}^{1 \times p}$)
 - ▶ Matrices have shapes (input dimension, output dimension)
 - ▶ Let $\Theta \in \mathbb{R}^{n \times p}$, $\mathbf{x} \in \mathbb{R}^{1 \times n}$, and $\mathbf{z} \in \mathbb{R}^{1 \times p}$
 - ▶ $\mathbf{z} = \mathbf{x} \Theta$ (or $\mathbf{x} \mathbf{W} + \mathbf{b}$)
 - ▶ More intuitive (input \rightarrow output)
 - ▶ Aligns with the convention in (mini)batch training that the first dimension is the batch size ("feature vectors are stacked row-wise")

General Advice

- ▶ Know your shapes!



Source

Multi-Class Classification

- ▶ What is the classification function g ?
- ▶ Let $\mathbf{z} = [z_1 \ \dots \ z_p]$ be a vector of scores for each class

Multi-Class Classification

- ▶ What is the classification function g ?
- ▶ Let $\mathbf{z} = [z_1 \ \dots \ z_p]$ be a vector of scores for each class
- ▶ Logistic regression: softmax function
 - ▶ $\text{softmax}(z_c) = \frac{e^{z_c}}{\sum_{k=1}^p e^{z_k}} = P(y = c|\mathbf{x}) = \hat{y}_c$
 - ▶ $\text{softmax}(\mathbf{z}) = \hat{\mathbf{y}}$ is a vector of probabilities for each class

Multi-Class Classification

- ▶ What is the classification function g ?
- ▶ Let $\mathbf{z} = [z_1 \ \dots \ z_p]$ be a vector of scores for each class
- ▶ Logistic regression: softmax function
 - ▶ $\text{softmax}(z_c) = \frac{e^{z_c}}{\sum_{k=1}^p e^{z_k}} = P(y = c|\mathbf{x}) = \hat{y}_c$
 - ▶ $\text{softmax}(\mathbf{z}) = \hat{\mathbf{y}}$ is a vector of probabilities for each class
- ▶ Perceptron: argmax function
 - ▶ $\text{argmax}_{k=1}^p(z_k) = \hat{y}$

Multi-Class Classification

- ▶ What is the classification function g ?
- ▶ Let $\mathbf{z} = [z_1 \ \dots \ z_p]$ be a vector of scores for each class
- ▶ Logistic regression: softmax function
 - ▶ $\text{softmax}(z_c) = \frac{e^{z_c}}{\sum_{k=1}^p e^{z_k}} = P(y = c|\mathbf{x}) = \hat{y}_c$
 - ▶ $\text{softmax}(\mathbf{z}) = \hat{\mathbf{y}}$ is a vector of probabilities for each class
- ▶ Perceptron: argmax function
 - ▶ $\text{argmax}_{k=1}^p(z_k) = \hat{y}$
 - ▶ What if there is a tie?
 - ▶ Do whatever `numpy.argmax` does

Naïve Bayes as a Linear Classifier

- ▶ Suppose we observe a document d . What is the most likely class \hat{c} ?

Naïve Bayes as a Linear Classifier

- ▶ Suppose we observe a document d . What is the most likely class \hat{c} ?

- ▶
$$P(c|d) = \frac{P(d|c)P(c)}{P(d)}$$

- ▶ Bayes' Rule

- ▶
$$\hat{c} = \operatorname{argmax}_{c \in C} P(d|c)P(c)$$

- ▶ $P(d)$ is the same for each class

- ▶
$$\hat{c} = \operatorname{argmax}_{c \in C} P(c) \prod_{i \in \text{positions}} P(w_i|c)$$

- ▶ Bag of words assumption, Naïve Bayes assumption

- ▶
$$\hat{c} = \operatorname{argmax}_{c \in C} \log P(c) + \sum_{i \in \text{positions}} \log P(w_i|c)$$

- ▶ If $xy = z$, then $\log(x) + \log(y) = \log(z)$

Naïve Bayes as a Linear Classifier

- ▶ $\hat{c} = \operatorname{argmax}_{c \in C} \sum_{w \in |V|} \left[(\text{count}(w, d)) (\log P(w|c)) \right] + \log P(c)$
- ▶ $\sum_{i \in \text{positions}} \log P(w_i|c) = \sum_{w \in |V|} \left[(\text{count}(w, d)) (\log P(w|c)) \right]$

Naïve Bayes as a Linear Classifier

- ▶ $\hat{c} = \operatorname{argmax}_{c \in C} \sum_{w \in |V|} \left[(\text{count}(w, d)) (\log P(w|c)) \right] + \log P(c)$
 - ▶ $\sum_{i \in \text{positions}} \log P(w_i|c) = \sum_{w \in |V|} \left[(\text{count}(w, d)) (\log P(w|c)) \right]$
- ▶ Let $x_w = \text{count}(w, d)$, $\ell_{cw} = \log P(w|c)$, and $p_c = \log P(c)$

$$\begin{aligned}\hat{c} &= \operatorname{argmax}_{c \in C} \sum_{w \in |V|} x_w \ell_{cw} + p_c \\ &= \operatorname{argmax}_{c \in C} (\mathbf{x} \cdot \boldsymbol{\ell}_c + p_c)\end{aligned}$$

Naïve Bayes as a Linear Classifier

- ▶ Let \mathbf{x} be a feature vector, $\mathcal{L} = \text{self.likelihood}$, and $\mathbf{p} = \text{self.prior}$

$$\mathbf{z} = \mathbf{x}\mathcal{L} + \mathbf{p}$$

$$\hat{c} = \underset{c \in C}{\operatorname{argmax}}(z_c)$$