

Structured Perceptrons

CS114B Lab 7

Kenneth Lai

March 10, 2023

Sequence Labeling

- ▶ Suppose we observe a list of words X . What are the respective parts of speech Y ?
 - ▶ What is $P(Y|X)$?

Hidden Markov Models

- ▶ Generative approach: Hidden Markov Models
 - ▶ $P(Y|X) \propto P(X, Y) = P(X|Y)P(Y)$

Hidden Markov Models

- ▶ Generative approach: Hidden Markov Models
 - ▶ $P(Y|X) \propto P(X, Y) = P(X|Y)P(Y)$
- ▶ Independence Assumptions
 - ▶ (First-order) Markov Assumption: the probability of a tag depends only on the previous tag
 - ▶
$$P(Y) = \prod_{i=1}^T P(y_i|y_{i-1})$$

Hidden Markov Models

- ▶ Generative approach: Hidden Markov Models

- ▶ $P(Y|X) \propto P(X, Y) = P(X|Y)P(Y)$

- ▶ Independence Assumptions

- ▶ (First-order) Markov Assumption: the probability of a tag depends only on the previous tag

- ▶
$$P(Y) = \prod_{i=1}^T P(y_i|y_{i-1})$$

- ▶ Output Independence: the probability of a word at time i depends only on the tag at time i

- ▶
$$P(X|Y) = \prod_{i=1}^T P(x_i|y_i)$$

Hidden Markov Models

$$\begin{aligned}P(Y|X) &\propto \prod_{i=1}^T P(x_i|y_i) \times \prod_{i=1}^T P(y_i|y_{i-1}) \\&\propto \prod_{i=1}^T \left(P(x_i|y_i) \times P(y_i|y_{i-1}) \right) \\ \log P(Y|X) &\propto \sum_{i=1}^T \left(\log P(x_i|y_i) + \log P(y_i|y_{i-1}) \right)\end{aligned}$$

Hidden Markov Models

$$\begin{aligned} P(Y|X) &\propto \prod_{i=1}^T P(x_i|y_i) \times \prod_{i=1}^T P(y_i|y_{i-1}) \\ &\propto \prod_{i=1}^T \left(P(x_i|y_i) \times P(y_i|y_{i-1}) \right) \\ \log P(Y|X) &\propto \sum_{i=1}^T \left(\log P(x_i|y_i) + \log P(y_i|y_{i-1}) \right) \end{aligned}$$

- ▶ In other words, $(\log) P(Y|X)$ decomposes into a product (or sum) of **local** parts
- ▶ This allows us to use dynamic programming

Viterbi Algorithm

- ▶ Not just for HMMs!
- ▶ Discriminative approaches:
 - ▶ Conditional random fields
 - ▶ Structured perceptrons

Viterbi Algorithm

- ▶ Not just for HMMs!
- ▶ Discriminative approaches:
 - ▶ Conditional random fields
 - ▶ Structured perceptrons
- ▶ As long as the “score” decomposes into a sum of local parts, we can use the Viterbi algorithm

Viterbi Algorithm

function VITERBI(*observations* of len T , *state-graph* of len N) **returns** *best-path*, *path-prob*

create a path probability matrix *viterbi*[N, T]

for each state s **from** 1 **to** N **do** ; initialization step

$viterbi[s, 1] \leftarrow \pi_s * b_s(o_1)$

$backpointer[s, 1] \leftarrow 0$

for each time step t **from** 2 **to** T **do** ; recursion step

for each state s **from** 1 **to** N **do**

$viterbi[s, t] \leftarrow \max_{s'=1}^N viterbi[s', t-1] * a_{s', s} * b_s(o_t)$

$backpointer[s, t] \leftarrow \operatorname{argmax}_{s'=1}^N viterbi[s', t-1] * a_{s', s} * b_s(o_t)$

$bestpathprob \leftarrow \max_{s=1}^N viterbi[s, T]$; termination step

$bestpathpointer \leftarrow \operatorname{argmax}_{s=1}^N viterbi[s, T]$; termination step

bestpath \leftarrow the path starting at state *bestpathpointer*, that follows *backpointer*[] to states back in time

return *bestpath*, *bestpathprob*

Figure 8.10 Viterbi algorithm for finding the optimal sequence of tags. Given an observation sequence and an HMM $\lambda = (A, B)$, the algorithm returns the state path through the HMM that assigns maximum likelihood to the observation sequence.

Viterbi Algorithm

- ▶ Input: observations of length T
- ▶ Also let N be the number of states and V be the size of the vocabulary

Viterbi Algorithm

- ▶ Input: observations of length T
- ▶ Also let N be the number of states and V be the size of the vocabulary
- ▶ Three Numpy arrays:
 - ▶ `self.initial` (π : shape $(N,)$)
 - ▶ `self.transition` (**A**: shape (N, N))
 - ▶ `self.emission` (**B**: shape (V, N))
- ▶ We assume you know how to fill them in
 - ▶ HMM: CS114A
 - ▶ Structured perceptron: in a few slides

Viterbi Algorithm

- ▶ Input: observations of length T
- ▶ Also let N be the number of states and V be the size of the vocabulary
- ▶ Three Numpy arrays:
 - ▶ `self.initial` (π : shape $(N,)$)
 - ▶ `self.transition` (**A**: shape (N, N))
 - ▶ `self.emission` (**B**: shape (V, N))
- ▶ We assume you know how to fill them in
 - ▶ HMM: CS114A
 - ▶ Structured perceptron: in a few slides
- ▶ For each sentence, create two Numpy arrays: (both with shape (N, T))
 - ▶ `v` (for viterbi)
 - ▶ `backpointer`

Initialization Step

- ▶ For each state s from 1 to N do
$$viterbi[s, 1] \leftarrow \pi_s + b_s(o_1)$$
- ▶ Note the use of $+$ instead of \times
- ▶ We'll see how to do this without the for loop in a bit

Recursion Step

- ▶ For each time step t from 2 to T do
For each state s from 1 to N do
$$viterbi[s, t] \leftarrow \max_{s'=1}^N viterbi[s', t-1] + a_{s',s} + b_s(o_t)$$
- ▶ Note the use of $+$ instead of \times
- ▶ We'll see how to do this without the (inner) for loop in a bit

Termination Step

- ▶ Best last tag is the argmax of the last column of v
- ▶ Follow backpointers in `backpointer`
 - ▶ Nothing fancy; we'll let you figure it out on your own
 - ▶ For HW4, you do not have to return the path (log-)probability/score, just the backtrace path

Structured Perceptrons

- ▶ Perceptrons for sequence labeling
 - ▶ Given a sequence $X = [x_1, \dots, x_T]$, predict labels $Y = [y_1, \dots, y_T]$
 - ▶ We do not care about the probability $P(Y|X)$, just which Y has the highest score Z

Structured Perceptrons

- ▶ Perceptrons for sequence labeling
 - ▶ Given a sequence $X = [x_1, \dots, x_T]$, predict labels $Y = [y_1, \dots, y_T]$
 - ▶ We do not care about the probability $P(Y|X)$, just which Y has the highest score Z
- ▶ $\hat{Y} = \operatorname{argmax}_{k \in K^T} Z_k$

Structured Perceptrons

- ▶ Score Z decomposes into a sum of local parts
 - ▶ At each time step i , for each possible combination of current tag y_i and previous tag y_{i-1} , compute a local score $z(y_i, y_{i-1})$

Structured Perceptrons

- ▶ Score Z decomposes into a sum of local parts
 - ▶ At each time step i , for each possible combination of current tag y_i and previous tag y_{i-1} , compute a local score $z(y_i, y_{i-1})$
 - ▶ Use the Viterbi algorithm to combine the local scores across the sequence, and find the argmax

Structured Perceptrons

- ▶ Suppose that at each time step i , we want to predict the current tag y_i using the following features:
 - ▶ Previous tag y_{i-1}
 - ▶ At the beginning of the sentence, let the previous tag y_{i-1} be the start symbol $\langle S \rangle$
 - ▶ Current word x_i

Structured Perceptrons

- ▶ Suppose that at each time step i , we want to predict the current tag y_i using the following features:
 - ▶ Previous tag y_{i-1}
 - ▶ At the beginning of the sentence, let the previous tag y_{i-1} be the start symbol $\langle S \rangle$
 - ▶ Current word x_i
- ▶ For simplicity, we will assume that these are the only features, and we will ignore the bias term

Structured Perceptrons

- ▶ Suppose that at each time step i , we want to predict the current tag y_i using the following features:
 - ▶ Previous tag y_{i-1}
 - ▶ At the beginning of the sentence, let the previous tag y_{i-1} be the start symbol $\langle S \rangle$
 - ▶ Current word x_i
- ▶ For simplicity, we will assume that these are the only features, and we will ignore the bias term
- ▶ Let $\mathbf{f}(X, y_i, y_{i-1}, i)$ be the feature vector at time step i
 - ▶ Using \mathbf{f} instead of \mathbf{x} , because features can include more than just the input

Structured Perceptrons

- ▶ We can arrange our weight matrix Θ as follows:

$$\begin{bmatrix} \text{---} \\ \text{---} \end{bmatrix}$$

Structured Perceptrons

- ▶ We can arrange our weight matrix Θ as follows:

$$\begin{bmatrix} \pi \\ \hline \\ \hline \end{bmatrix}$$

- ▶ Initial features

- ▶ $y_{i-1} = \langle S \rangle, y_i = \dots$

Structured Perceptrons

- ▶ We can arrange our weight matrix Θ as follows:

$$\begin{bmatrix} \pi \\ \hline \mathbf{A} \\ \hline \end{bmatrix}$$

- ▶ Initial features
 - ▶ $y_{i-1} = \langle S \rangle, y_i = \dots$
- ▶ Transition features
 - ▶ $y_{i-1} = \dots, y_i = \dots$

Structured Perceptrons

- ▶ We can arrange our weight matrix Θ as follows:

$$\begin{bmatrix} \pi \\ \hline \mathbf{A} \\ \hline \mathbf{B} \end{bmatrix}$$

- ▶ Initial features
 - ▶ $y_{i-1} = \langle S \rangle, y_i = \dots$
- ▶ Transition features
 - ▶ $y_{i-1} = \dots, y_i = \dots$
- ▶ Emission features
 - ▶ $x_i = \dots, y_i = \dots$

Initialization Step

- ▶ We want to compute local scores $z(y_1, \langle S \rangle)$ for each possible y_1
 - ▶ These are the elements of $\mathbf{z}_1 = \mathbf{f}(X, y_1, \langle S \rangle, 1) \cdot \Theta$

Initialization Step

$$\mathbf{z}_1 = \mathbf{f}(X, y_1, \langle S \rangle, 1) \cdot \Theta$$

$$= \begin{bmatrix} & | & & | \end{bmatrix} \cdot \begin{bmatrix} \pi \\ \hline \mathbf{A} \\ \hline \mathbf{B} \end{bmatrix}$$

Initialization Step

$$\mathbf{z}_1 = \mathbf{f}(X, y_1, \langle S \rangle, 1) \cdot \Theta$$

$$= \left[\begin{array}{c|c} 1 & \end{array} \right]$$

$$\cdot \begin{bmatrix} \pi \\ \hline \mathbf{A} \\ \hline \mathbf{B} \end{bmatrix}$$

- We know that $y_{i-1} = \langle S \rangle$

Initialization Step

$$\mathbf{z}_1 = \mathbf{f}(X, y_1, \langle S \rangle, 1) \cdot \Theta$$

$$= \left[\begin{array}{c|c|c} 1 & \mathbf{0} & \end{array} \right] \cdot \left[\begin{array}{c} \hline \pi \\ \hline \mathbf{A} \\ \hline \mathbf{B} \\ \hline \end{array} \right]$$

- We know that y_{i-1} cannot be any other tag

Initialization Step

$$\mathbf{z}_1 = \mathbf{f}(X, y_1, \langle S \rangle, 1) \cdot \Theta$$

$$= \left[\begin{array}{c|c|c} 1 & \mathbf{0} & \mathbf{1}_{\{x_1 = o_1\}} \end{array} \right] \cdot \left[\begin{array}{c} \pi \\ \hline \mathbf{A} \\ \hline \mathbf{B} \end{array} \right]$$

- One-hot vector of the first word

Initialization Step

$$\mathbf{z}_1 = \mathbf{f}(X, y_1, \langle S \rangle, 1) \cdot \Theta$$

$$= \begin{bmatrix} 1 & \mathbf{0} & \mathbf{1}_{\{x_1 = o_1\}} \end{bmatrix} \cdot \begin{bmatrix} \pi \\ \hline \mathbf{A} \\ \hline \mathbf{B} \end{bmatrix}$$

$$= 1 \cdot \pi + \mathbf{0} \cdot \mathbf{A} + \mathbf{1}_{\{x_1 = o_1\}} \cdot \mathbf{B}$$

$$= \pi + \mathbf{B}[o_1]$$

Initialization Step

$$\mathbf{z}_1 = \mathbf{f}(X, y_1, \langle S \rangle, 1) \cdot \Theta$$

$$= \begin{bmatrix} 1 & | & \mathbf{0} & | & \mathbf{1}\{x_1 = o_1\} \end{bmatrix} \cdot \begin{bmatrix} \pi \\ \hline \mathbf{A} \\ \hline \mathbf{B} \end{bmatrix}$$

$$= 1 \cdot \pi + \mathbf{0} \cdot \mathbf{A} + \mathbf{1}\{x_1 = o_1\} \cdot \mathbf{B}$$

$$= \pi + \mathbf{B}[o_1]$$

- These local scores go into the first column of the Viterbi trellis

Recursion Step

- ▶ We want to compute local scores $z(y_i, y_{i-1})$ for each possible combination of y_i and y_{i-1}
 - ▶ We can stack the feature vectors for each possible y_{i-1} , $\mathbf{f}(X, y_i, y_{i-1}, i)$, on top of each other, in order
 - ▶ Form a feature matrix \mathbf{F}_i

Recursion Step

- ▶ We want to compute local scores $z(y_i, y_{i-1})$ for each possible combination of y_i and y_{i-1}
 - ▶ We can stack the feature vectors for each possible y_{i-1} , $\mathbf{f}(X, y_i, y_{i-1}, i)$, on top of each other, in order
 - ▶ Form a feature matrix \mathbf{F}_i
 - ▶ Compute $\mathbf{Z}_i = \mathbf{F}_i \cdot \Theta$

Recursion Step

$$\mathbf{Z}_i = \mathbf{F}_i \cdot \boldsymbol{\Theta}$$

$$= \begin{bmatrix} | & | & | \end{bmatrix} \cdot \begin{bmatrix} \pi \\ \hline \mathbf{A} \\ \hline \mathbf{B} \end{bmatrix}$$

Recursion Step

$$\mathbf{Z}_i = \mathbf{F}_i \cdot \Theta$$

$$= \left[\begin{array}{c|c} \mathbf{0} & \end{array} \right] \cdot \left[\begin{array}{c} \pi \\ \hline \mathbf{A} \\ \hline \mathbf{B} \end{array} \right]$$

- We know that $y_{i-1} \neq \langle S \rangle$

Recursion Step

$$\mathbf{Z}_i = \mathbf{F}_i \cdot \Theta$$

$$= \left[\begin{array}{c|c} \mathbf{0} & \mathbf{I} \end{array} \right] \cdot \left[\begin{array}{c} \pi \\ \hline \mathbf{A} \\ \hline \mathbf{B} \end{array} \right]$$

► Identity matrix!

Recursion Step

$$\mathbf{Z}_i = \mathbf{F}_i \cdot \Theta$$

$$= \left[\begin{array}{c|c|c} \mathbf{0} & \mathbf{I} & \mathbf{1}_{\{x_i = o_i\}} \end{array} \right] \cdot \left[\begin{array}{c} \pi \\ \hline \mathbf{A} \\ \hline \mathbf{B} \end{array} \right]$$

- Stack of one-hot vectors

Recursion Step

$$\mathbf{Z}_i = \mathbf{F}_i \cdot \Theta$$

$$= \left[\begin{array}{c|c|c} \mathbf{0} & \mathbf{I} & \mathbf{1}_{\{x_i = o_i\}} \end{array} \right] \cdot \left[\begin{array}{c} \pi \\ \hline \mathbf{A} \\ \hline \mathbf{B} \end{array} \right]$$

$$= \mathbf{0} \cdot \pi + \mathbf{I} \cdot \mathbf{A} + \mathbf{1}_{\{x_i = o_i\}} \cdot \mathbf{B}$$

$$= \mathbf{A} + \mathbf{B}[o_i]$$

Recursion Step

$$\mathbf{Z}_i = \mathbf{F}_i \cdot \Theta$$

$$= \left[\begin{array}{c|c|c} \mathbf{0} & \mathbf{I} & \mathbf{1}_{\{x_i = o_i\}} \end{array} \right] \cdot \left[\begin{array}{c} \pi \\ \hline \mathbf{A} \\ \hline \mathbf{B} \end{array} \right]$$

$$= \mathbf{0} \cdot \pi + \mathbf{I} \cdot \mathbf{A} + \mathbf{1}_{\{x_i = o_i\}} \cdot \mathbf{B}$$

$$= \mathbf{A} + \mathbf{B}[o_i]$$

- Use the Viterbi algorithm to combine these local scores with scores from the rest of the sequence

Recursion Step

- ▶ Use the Viterbi algorithm to combine these local scores with scores from the rest of the sequence

$$\mathbf{Z}_i = \mathbf{F}_i \cdot \mathbf{\Theta} = \mathbf{A} + \mathbf{B}[o_i]$$

Recursion Step

- Use the Viterbi algorithm to combine these local scores with scores from the rest of the sequence

$$\mathbf{Z}_i = \mathbf{F}_i \cdot \mathbf{\Theta} = \mathbf{A} + \mathbf{B}[o_i]$$

$$\text{Viterbi}[:, i] = \max(\text{Viterbi}[:, i - 1 : i] + \mathbf{A} + \mathbf{B}[o_i], \text{axis}=0)$$

Termination Step

- ▶ Best last tag is the argmax of the last column of v
- ▶ Follow backpointers in `backpointer`
 - ▶ Nothing fancy; we'll let you figure it out on your own
 - ▶ For HW4, you do not have to return the path (log-)probability/score, just the backtrace path

Structured Perceptron Learning Algorithm

- ▶ Use the Viterbi algorithm to compute the best tag sequence

Structured Perceptron Learning Algorithm

- ▶ Use the Viterbi algorithm to compute the best tag sequence
- ▶ If $\hat{Y} = Y$, then do nothing

Structured Perceptron Learning Algorithm

- ▶ Use the Viterbi algorithm to compute the best tag sequence
- ▶ If $\hat{Y} = Y$, then do nothing
- ▶ Else:
 - ▶ Increment weights for features in Y , decrement weights for features in \hat{Y}

Structured Perceptron Learning Algorithm

- ▶ Use the Viterbi algorithm to compute the best tag sequence
- ▶ If $\hat{Y} = Y$, then do nothing
- ▶ Else:
 - ▶ Increment weights for features in Y , decrement weights for features in \hat{Y}
 - ▶ In other words, for each time step i :
 - ▶ Increment weights for features in y_i , decrement weights for features in \hat{y}_i

Structured Perceptron Learning Algorithm

- ▶ Use the Viterbi algorithm to compute the best tag sequence
- ▶ If $\hat{Y} = Y$, then do nothing
- ▶ Else:
 - ▶ Increment weights for features in Y , decrement weights for features in \hat{Y}
 - ▶ In other words, for each time step i :
 - ▶ Increment weights for features in y_i , decrement weights for features in \hat{y}_i
 - ▶ Nothing fancy; no Numpy tricks needed