

# Recurrent Neural Networks and Transformers

## CS114B Lab 11

Kenneth Lai

April 14, 2023

# Sequence Labeling

- ▶ Suppose we observe a list of words  $X$ . What are the respective parts of speech  $Y$ ?
  - ▶ What is  $P(Y|X)$ ?

# Sequence Labeling

- ▶ Suppose we observe a list of words  $X$ . What are the respective parts of speech  $Y$ ?
  - ▶ What is  $P(Y|X)$ ?
- ▶ Generative approach: Hidden Markov Models

# Sequence Labeling

- ▶ Suppose we observe a list of words  $X$ . What are the respective parts of speech  $Y$ ?
  - ▶ What is  $P(Y|X)$ ?
- ▶ Generative approach: Hidden Markov Models
- ▶ Discriminative approaches:
  - ▶ Conditional random fields
  - ▶ Structured perceptrons

# Sequence Labeling

- ▶ Suppose we observe a list of words  $X$ . What are the respective parts of speech  $Y$ ?
  - ▶ What is  $P(Y|X)$ ?
- ▶ Generative approach: Hidden Markov Models
- ▶ Discriminative approaches:
  - ▶ Conditional random fields
  - ▶ Structured perceptrons
  - ▶ Neural networks


# Sequence Labeling

- ▶ Discriminative approaches:
  - ▶ At each time step, use local features to compute local scores, and use the Viterbi algorithm to make predictions for the whole sentence
    - ▶ Conditional random fields
    - ▶ Structured perceptrons

# Sequence Labeling

- ▶ Discriminative approaches:
  - ▶ At each time step, use local features to compute local scores, and use the Viterbi algorithm to make predictions for the whole sentence
    - ▶ Conditional random fields
    - ▶ Structured perceptrons
  - ▶ Use features from other time steps, but make independent predictions at each time step
    - ▶ Neural networks


# Feedforward Neural Networks



- ▶ Output layer
- ▶ Hidden layer(s)
- ▶ Input layer
- ▶  $h = g(xW)$
- ▶  $\hat{y} = g'(hV)$




# Feedforward Neural Networks




- ▶ Output layer
- ▶ Hidden layer(s)
- ▶ Input layer
- ▶  $\mathbf{h} = g(\mathbf{x}\mathbf{W})$
- ▶  $\hat{\mathbf{y}} = g'(\mathbf{h}\mathbf{V})$ 
  - ▶ We will assume that the dummy feature 1 is part of  $\mathbf{x}$  and  $\mathbf{h}$ , and that the bias term is part of  $\mathbf{W}$  and  $\mathbf{V}$ , etc.

# Neural Networks for Sequence Labeling




# Neural Networks for Sequence Labeling




- ▶ At each time  $i$ , the input to the neural network consists of:
  - ▶ Current word (or other input) vector  $x_i$
  - ▶ History/(past) context vector  $h_{i-1}$

# Neural Networks for Sequence Labeling




- ▶ At each time  $i$ , the input to the neural network consists of:
  - ▶ Current word (or other input) vector  $\mathbf{x}_i$
  - ▶ History/(past) context vector  $\mathbf{h}_{i-1}$
- ▶  $\mathbf{h}_i = g(\mathbf{x}_i \mathbf{W} + \mathbf{h}_{i-1} \mathbf{U})$

# Neural Networks for Sequence Labeling




# Neural Networks for Sequence Labeling




- The output of the hidden state at one time step is the history/past context input for the next time step!

# Neural Networks for Sequence Labeling



- ▶ The output of the hidden state at one time step is the history/past context input for the next time step!
- ▶ What context information is embedded in  $\mathbf{h}_{i-1}$ ?
  - ▶ Previous word  $\mathbf{x}_{i-1}$
  - ▶ Previous context  $\mathbf{h}_{i-2}$


# Neural Networks for Sequence Labeling



- ▶ The output of the hidden state at one time step is the history/past context input for the next time step!
- ▶ What context information is embedded in  $\mathbf{h}_{i-1}$ ?
  - ▶ Previous word  $\mathbf{x}_{i-1}$
  - ▶ Previous context  $\mathbf{h}_{i-2}$ 
    - ▶ Previous previous word  $\mathbf{x}_{i-2}$
    - ▶ Previous previous context  $\mathbf{h}_{i-3}$




# Neural Networks for Sequence Labeling




- ▶ The output of the hidden state at one time step is the history/past context input for the next time step!
- ▶ What context information is embedded in  $h_{i-1}$ ?
  - ▶ All previous words

# Neural Networks for Sequence Labeling




- ▶ The output of the hidden state at one time step is the history/past context input for the next time step!
- ▶ What context information is embedded in  $\mathbf{h}_{i-1}$ ?
  - ▶ All previous words
  - ▶ What about previous parts of speech (as in HMMs, CRFs, structured perceptrons)?
    - ▶ To learn more, take CS231A in the fall!

# Recurrent Neural Networks




# Recurrent Neural Networks




- Neural networks in which the output of a layer in one time step is input to a layer in the next time step

# RNN Language Models




# RNN Language Models




- ▶ Sequence labeling: predict current tag given current word, history
- ▶ Language modeling: predict next word given current word, history

# RNN Language Models




- ▶ Sequence labeling: predict current tag given current word, context
- ▶ Language modeling: predict next word given current word, context

# RNNs for Text Classification






# RNNs for Text Classification




- ▶ What context information is embedded in  $\mathbf{h}_T$ ?
  - ▶ Current word  $\mathbf{x}_T$
  - ▶ Context  $\mathbf{h}_{T-1}$

# RNNs for Text Classification



- ▶ What context information is embedded in  $\mathbf{h}_T$ ?
  - ▶ All words (i.e. the whole text)
- ▶ Use  $\mathbf{h}_T$  to predict class  $\hat{\mathbf{y}}_T$  of entire document

# RNNs for Text Classification




- ▶ What context information is embedded in  $\mathbf{h}_T$ ?
  - ▶ All words (i.e. the whole text)
- ▶ Use  $\mathbf{h}_T$  to predict class  $\hat{\mathbf{y}}_T$  of entire document
  - ▶ Ignore other outputs


# Backpropagation

- ▶ For each matrix of weights  $\mathbf{W}$ , starting from the output and working backwards:
  - ▶ Compute gradient  $(\nabla L)^{[W]}$
- ▶ For each matrix of weights  $\mathbf{W}$ :
  - ▶ Move in direction of negative gradient

# Backpropagation




# Backpropagation




- Compute gradient  $(\nabla L)^{[V]}$

# Backpropagation




- ▶ Compute gradient  $(\nabla L)^{[V]}$
- ▶ Use  $(\nabla L)^{[V]}$  to compute gradient  $(\nabla L)^{[W]}$

# Backpropagation Through Time






# Backpropagation Through Time




- ▶ Start at the end of the text and work backwards
  - ▶ Let  $(\nabla L)_{i,j}^{[\mathbf{W}]}$  denote the part of the gradient for weight matrix  $\mathbf{W}$  at time  $i$  that comes from the output at time  $j$

# Backpropagation Through Time




- ▶ Start at the end of the text and work backwards
  - ▶ Compute gradient  $(\nabla L)_{T,T}^{[V]}$

# Backpropagation Through Time




- ▶ Start at the end of the text and work backwards
  - ▶ Compute gradient  $(\nabla L)_{T,T}^{[V]}$
  - ▶ Use  $(\nabla L)_{T,T}^{[V]}$  to compute gradients  $(\nabla L)_{T,T}^{[W]}$  and  $(\nabla L)_{T,T}^{[U]}$

# Backpropagation Through Time




- ▶ Start at the end of the text and work backwards
  - ▶ Compute gradient  $(\nabla L)_{T,T}^{[V]}$
  - ▶ Use  $(\nabla L)_{T,T}^{[V]}$  to compute gradients  $(\nabla L)_{T,T}^{[W]}$  and  $(\nabla L)_{T,T}^{[U]}$
  - ▶ Use  $(\nabla L)_{T,T}^{[U]}$  to compute gradients  $(\nabla L)_{T-1,T}^{[W]}$  and  $(\nabla L)_{T-1,T}^{[U]}$
  - ▶ etc.

# Backpropagation Through Time



- ▶ Start at the end of the text and work backwards
  - ▶ Compute gradient  $(\nabla L)_{T-1, T-1}^{[V]}$

# Backpropagation Through Time



- ▶ Start at the end of the text and work backwards
  - ▶ Compute gradient  $(\nabla L)_{T-1, T-1}^{[V]}$
  - ▶ Use  $(\nabla L)_{T-1, T-1}^{[V]}$  to compute gradients  $(\nabla L)_{T-1, T-1}^{[W]}$  and  $(\nabla L)_{T-1, T-1}^{[U]}$
  - ▶ etc.

# Backpropagation Through Time

- ▶ The overall gradient for a weight matrix  $\mathbf{W}$  is the sum of the gradients at each time  $i$  from each output  $\hat{\mathbf{y}}_j$

- ▶ 
$$(\nabla L)^{[\mathbf{W}]} = \sum_{j=1}^T \sum_{i=1}^j (\nabla L)_{i,j}^{[\mathbf{W}]}$$

# Backpropagation Through Time


- ▶ The overall gradient for a weight matrix  $\mathbf{W}$  is the sum of the gradients at each time  $i$  from each output  $\hat{\mathbf{y}}_j$

- ▶  $(\nabla L)^{[\mathbf{W}]} = \sum_{j=1}^T \sum_{i=1}^j (\nabla L)_{i,j}^{[\mathbf{W}]}$


- ▶ Then move in direction of negative gradient



# Recurrent Neural Networks



# Recurrent Neural Networks




- ▶ Output  $\hat{y}_i$  depends on hidden state  $h_i$  (i.e. current word  $x_i$  and history/(past) context  $h_{i-1}$ )
- ▶ What about future context?


# Bidirectional RNNs

- ▶ Idea: Train two RNNs: passing the input into one forward and one **backward**
- ▶ Output  $\hat{\mathbf{y}}_i$  depends on forward hidden state  $\vec{\mathbf{h}}_i$  and backward hidden state  $\overleftarrow{\mathbf{h}}_i$


# Forward RNN




# Backward RNN



# Bidirectional RNN




# Bidirectional RNNs for Text Classification



- ▶  $\vec{h}_T$  encodes the whole text
  - ▶ Use  $\vec{h}_T$  to predict class  $\hat{y}_T$  of entire document


# Bidirectional RNNs for Text Classification



- ▶  $\overrightarrow{\mathbf{h}}_T$  encodes the whole text
  - ▶ Use  $\overrightarrow{\mathbf{h}}_T$  to predict class  $\hat{\mathbf{y}}_T$  of entire document
- ▶  $\overleftarrow{\mathbf{h}}_1$  also encodes the whole text
  - ▶ Use  $\overleftarrow{\mathbf{h}}_1$  to predict class  $\hat{\mathbf{y}}_1$  of entire document



# Bidirectional RNNs for Text Classification



- Use  $\overrightarrow{\mathbf{h}}_T$  and  $\overleftarrow{\mathbf{h}}_1$  to predict class  $\hat{\mathbf{y}}$  of entire document

# Context and Long-Distance Dependencies

- ▶  $\mathbf{h}_{i-1}$  encodes the (past, in a forward RNN) context  $\mathbf{x}_1, \dots, \mathbf{x}_{i-1}$ 
  - ▶ But mostly  $\mathbf{x}_{i-1}$ , less  $\mathbf{x}_{i-2}$ , even less  $\mathbf{x}_{i-3}$ , ..., very little  $\mathbf{x}_1$
- ▶ Context is **local**

# Context and Long-Distance Dependencies

- ▶ Example: subject-verb agreement
- ▶ The flights the airline (was/were) cancelling (was/were) full.

# Context and Long-Distance Dependencies

- ▶ Example: subject-verb agreement
- ▶ The flights the **airline was** cancelling (was/were) full.
  - ▶ The context for “**was**” is mostly “**airline**”

# Context and Long-Distance Dependencies

- ▶ Example: subject-verb agreement
- ▶ The **flights** the **airline** **was** cancelling **were** full.
  - ▶ The context for “**was**” is mostly “**airline**”
  - ▶ The context for “**were**” is mostly “cancelling”, “**was**”, “**airline**”
    - ▶ Very little “**flights**”

# Context and Long-Distance Dependencies

- ▶ Two approaches to handling long-distance dependencies:
  - ▶ Memory-based (e.g. long short-term)



- ▶  does this

# Context and Long-Distance Dependencies

- ▶ Two approaches to handling long-distance dependencies:
  - ▶ Memory-based (e.g. long short-term)



- ▶ does this


- ▶ Attention-based

- ▶ At each time step, the model explicitly computes which other words to pay attention to




- ▶ does this

# Simple RNN






# Simple RNN




Source

# Long Short-Term Memory



Source


# Long Short-Term Memory



Source

- ▶ Separate memory (cell) state
  - ▶ Reading from and writing to memory controlled by **gates**
    - ▶ Each gate contains one or two neural network layers


# Long Short-Term Memory



Source

- ▶ Separate memory (cell) state
  - ▶ Reading from and writing to memory controlled by **gates**
    - ▶ Each gate contains one or two neural network layers
  - ▶ State **persists** across time
    - ▶ May remember information from long ago
    - ▶ Gradients for memory don't decay with time

# Forget Gate




$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

Source

- Neural network layer with logistic activation function

# Forget Gate




$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

Source

- ▶ Neural network layer with logistic activation function
- ▶ Element-wise multiplication of forget gate output with memory state
  - ▶ **Mask**: What parts of memory to forget/remember?

# Input Gate

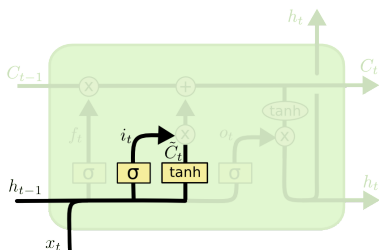


$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$
$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

Source

- ▶ Two parts
  1. Candidate choice
    - ▶ Logistic activation function
    - ▶ What parts of memory to update?

# Input Gate



$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$


$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

Source

- ▶ Two parts
  1. Candidate choice
    - ▶ Logistic activation function
    - ▶ What parts of memory to update?
  2. Candidate values
    - ▶ Tanh activation function
    - ▶ How much to update them by?



# Input Gate




$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

Source

- ▶ Element-wise multiplication of two outputs
- ▶ Then element-wise addition with memory state

# Output Gate




$$o_t = \sigma(W_o [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t * \tanh(C_t)$$

Source

- ▶ Logistic activation function
  - ▶ What parts of memory to output?

# Output Gate



$$o_t = \sigma(W_o [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t * \tanh(C_t)$$

Source

- ▶ Logistic activation function
  - ▶ What parts of memory to output?
- ▶ Element-wise multiplication with tanh of memory state
  - ▶ This is the “hidden layer output” that gets passed on to the output layer/next time step

# Context and Long-Distance Dependencies

- ▶ Two approaches to handling long-distance dependencies:
  - ▶ Memory-based (e.g. long short-term)



- ▶ does this

- ▶ Attention-based

- ▶ At each time step, the model explicitly computes which other words to pay attention to



- ▶ does this

# Attention

- ▶ Scaled dot-product self-attention

# Attention

- ▶ Scaled dot-product self-attention
  - ▶ Scaled dot-product: how to compute the relevance of the other words
  - ▶ Self-attention: paying attention to the input sequence itself (rather than some output sequence)

# Attention

- ▶ Input: pre-trained word vectors (e.g. from word2vec)

# Attention

- ▶ Input: pre-trained word vectors (e.g. from word2vec)
- 1. Compute **query**, **key**, and **value** vectors for each input vector
  - ▶ Matrix multiplication



# Attention

- ▶ For each word  $i$ :
  2. Compute the dot product of query  $q_i$  with key  $k_j$  for each word  $j$

# Attention

- ▶ For each word  $i$ :
  2. Compute the dot product of query  $q_i$  with key  $k_j$  for each word  $j$
  3. Scale the dot product (e.g. Vaswani et al. (2017) divide by 8)
    - ▶ Leads to more stable gradients


# Attention

- ▶ For each word  $i$ :
  2. Compute the dot product of query  $q_i$  with key  $k_j$  for each word  $j$
  3. Scale the dot product (e.g. Vaswani et al. (2017) divide by 8)
    - ▶ Leads to more stable gradients
  4. Softmax
    - ▶ Numbers  $\rightarrow$  probabilities (or more generally, numbers between 0 and 1, that add up to 1)

# Attention

- ▶ For each word  $i$ :
  2. Compute the dot product of query  $q_i$  with key  $k_j$  for each word  $j$
  3. Scale the dot product (e.g. Vaswani et al. (2017) divide by 8)
    - ▶ Leads to more stable gradients
  4. Softmax
    - ▶ Numbers  $\rightarrow$  probabilities (or more generally, numbers between 0 and 1, that add up to 1)
  5. Compute the weighted sum of values  $v_j$  for each word  $j$ 
    - ▶ Weights = softmax output from previous step

# Attention



# Attention

- ▶ Output: weighted sum of value vectors (modulo some more advanced topics)
  - ▶ Multi-head attention
  - ▶ Positional encodings
  - ▶ Residual connections
  - ▶ Layer normalization

# Attention

- ▶ Output: weighted sum of value vectors (modulo some more advanced topics)
  - ▶ Multi-head attention
  - ▶ Positional encodings
  - ▶ Residual connections
  - ▶ Layer normalization
- ▶ See Jay Alammar's [The Illustrated Transformer](#) for more details!

# Transformers

- ▶ “Attention Is All You Need” (Vaswani et al. 2017)
- ▶ No recurrence, relies entirely on attention (and feedforward layers) to capture global dependencies



# Transformers

- ▶ “Attention Is All You Need” (Vaswani et al. 2017)
- ▶ No recurrence, relies entirely on attention (and feedforward layers) to capture global dependencies
  - ▶ Recurrent neural networks are inherently sequential, processing one word at a time
  - ▶ Transformers are more parallel, looking at the entire sequence at once

# Transformers

- ▶ “Attention Is All You Need” (Vaswani et al. 2017)
- ▶ No recurrence, relies entirely on attention (and feedforward layers) to capture global dependencies
  - ▶ Recurrent neural networks are inherently sequential, processing one word at a time
  - ▶ Transformers are more parallel, looking at the entire sequence at once
    - ▶ More efficient, especially on GPUs
    - ▶ Also scores better on many NLP tasks

# Transformers

- ▶ Three “building blocks” of practical transformer models (Wolf et al. 2020):

# Transformers

- ▶ Three “building blocks” of practical transformer models (Wolf et al. 2020):
  1. Tokenizer
    - ▶ What is the input to the transformer?
    - ▶ Examples: words, WordPieces, characters, etc.

# Transformers

- ▶ Three “building blocks” of practical transformer models (Wolf et al. 2020):
  2. Transformer
    - ▶ The model itself (minus the output)

# Transformers

- ▶ Three “building blocks” of practical transformer models (Wolf et al. 2020):


## 2. Transformer

- ▶ The model itself (minus the output)



- ▶ : 12-24 encoder layers

- ▶ Encoder layer = (shared) attention layer + (individual) feedforward layers



# Transformers

- ▶ Three “building blocks” of practical transformer models (Wolf et al. 2020):
  3. Head
    - ▶ What is the output of the transformer?

# Transformers

- ▶ Three “building blocks” of practical transformer models (Wolf et al. 2020):

## 3. Head

- ▶ What is the output of the transformer?
- ▶ Models are **pre-trained** using one or more heads, but you can then **fine-tune** a model using a different head (task)



# Transformers

- ▶ Three “building blocks” of practical transformer models (Wolf et al. 2020):

## 3. Head

- ▶ What is the output of the transformer?
- ▶ Models are **pre-trained** using one or more heads, but you can then **fine-tune** a model using a different head (task)




- ▶ : Masked LM (Cloze) and NSP (Next Sentence Prediction)
  - ▶ Masked LM: mask 15% of input words at random, predict masked words
  - ▶ NSP: given sentences  $A$  and  $B$ , does  $B$  follow  $A$ ?

► Table from Wolf et al. (2020)

Name	Input	Heads		Ex. Datasets
		Output	Tasks	
Language Modeling	$x_{1:n-1}$	$x_n \in \mathcal{V}$	Generation	WikiText-103
Sequence Classification	$x_{1:L}$	$y \in \mathcal{C}$	Classification, Sentiment Analysis	GLUE, SST, MNLI
Question Answering	$x_{1:M}, x_{M+1:N}$	$y \text{ span } [1 : N]$	QA, Reading Comprehension	SQuAD, Natural Questions
Token Classification	$x_{1:L}$	$y_{1:L} \in \mathcal{C}^N$	NER, Tagging	OntoNotes, WNUT
Multiple Choice	$x_{1:L}, \mathcal{X}$	$y \in \mathcal{V}$	Text Selection	SWAG, ARC
Masked LM	$x_{1:L} \setminus \mathcal{N}$	$x_n \in \mathcal{V}$	Pretraining	WikiText, C4
Conditional Generation	$x_{1:L}$	$y_{1:M} \in \mathcal{V}^M$	Translation, Summarization	WMT, IWSLT, CNNDM, XSum

Transformers	
	Masked $[x_{1:N}, \dots, x_n]$
BERT	(Devlin et al., 2018)
RoBERTa	(Liu et al., 2019a)
	Autoregressive $[x_{1:n-1}, \dots, x_n]$
GPT / GPT-2	(Radford et al., 2019)
Trans-XL	(Dai et al., 2019)
XLNet	(Yang et al., 2019)
	Seq-to-Seq $[x_{1:N} \rightarrow x_{1:N}]$
BART	(Lewis et al., 2019)
T5	(Raffel et al., 2019)
MarianMT	(J. Dörmunt et al., 2018)
	Specialty: Multilingual
MMBT	(Kiehl et al., 2019)
	Specialty: Long-Distance
Reformer	(Kitaev et al., 2020)
Longformer	(Beltagy et al., 2020)
	Specialty: Efficient
ALBERT	(Lan et al., 2019)
Electra	(Clark et al., 2020)
DistilBERT	(Sanh et al., 2019)
	Specialty: Multilingual
XLM/RoBERTa	(Lample and Conneau, 2019b)



Tokenizers	
Name	Ex. Uses
Character-Level BPE	NMT, GPT
Byte-Level BPE	GPT-2
WordPiece	BERT
SentencePiece	XLNet
Unigram	LM
Character	Reformer
Custom	Bio-Chem

Figure 2: The *Transformers* library. **(Diagram-Right)** Each model is made up of a Tokenizer, *Transformer*, and Head. The model is pretrained with a fixed head and can then be further fine-tuned with alternate heads for different tasks. **(Bottom)** Each model uses a specific Tokenizer either implemented in Python or in Rust. These often differ in small details, but need to be in sync with preprocessing. **(Left)** Transformer architectures specialized for different tasks, e.g. understanding versus generation, or for specific use-cases, e.g. speed, image+text. **(Top)** heads allow a Transformer to be used for different tasks. Here we assume the input token sequence is  $x_{1:n}$  from a vocabulary  $V$ , and  $y$  represents different possible outputs, possibly from a class set  $C$ . Example datasets represent a small subset of example code distributed with the library.